

Two-Sample Test for Laws of Random Probabilities via Optimal Transport



George Kanchaveli

Bocconi University

The 5th Italian Meeting on Probability and Mathematical Statistics

June 11, 2026

Joint work with



Marta Catalano


LUISS 
Università di Roma




Hugo Lavenant




Two-Sample Testing in $\mathbb{R}$



Two-Sample Testing in $\mathbb{R}$



Two-Sample Testing in $\mathbb{R}$




Do they come from the same law?

Two-Sample Testing in $\mathbb{R}$

(unknown) $P, Q \in \mathcal{P}(\mathbb{R})$

(Observed) $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} P, \quad Y_1, \dots, Y_n \stackrel{\text{i.i.d.}}{\sim} Q.$

Test $H_0 : P = Q$ vs $H_1 : P \neq Q.$



Do they come from the same law?

Two-Sample Testing in $\mathbb{R}$


(unknown) $P, Q \in \mathcal{P}(\mathbb{R})$

(Observed) $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} P, \quad Y_1, \dots, Y_n \stackrel{\text{i.i.d.}}{\sim} Q.$

Test $H_0 : P = Q$ vs $H_1 : P \neq Q.$

Example: Kolmogorov–Smirnov Test

$$T_n = \sup_{x \in \mathbb{R}} |\hat{F}_n(x) - \hat{G}_n(x)|$$



Do they come from the same law?

Two-Sample Testing in $\mathbb{R}$


(unknown) $P, Q \in \mathcal{P}(\mathbb{R})$

(Observed) $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} P, \quad Y_1, \dots, Y_n \stackrel{\text{i.i.d.}}{\sim} Q.$

Test $H_0 : P = Q$ vs $H_1 : P \neq Q$.

Example: Kolmogorov–Smirnov Test

$$T_n = \sup_{x \in \mathbb{R}} |\hat{F}_n(x) - \hat{G}_n(x)|$$



Do they come from the same law?

Focus: *distance-based tests* — reject H_0 if samples are *far apart*.

Two-Sample Testing in $\mathbb{R}$


(unknown) $P, Q \in \mathcal{P}(\mathbb{R})$

(Observed) $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} P, \quad Y_1, \dots, Y_n \stackrel{\text{i.i.d.}}{\sim} Q.$

Test $H_0 : P = Q$ vs $H_1 : P \neq Q.$

Example: Kolmogorov–Smirnov Test

$$T_n = \sup_{x \in \mathbb{R}} |\hat{F}_n(x) - \hat{G}_n(x)|$$



Do they come from the same law?

Focus: *distance-based tests* — reject H_0 if samples are *far apart*.

$$d\left(\frac{1}{n} \sum_{i=1}^n \delta_{X_i}, \frac{1}{n} \sum_{i=1}^n \delta_{Y_i}\right) \text{ is large}$$

Two-Sample Testing in $\mathbb{R}^d$

Works in $\mathbb{R}^d$:

- Gretton et al. (2012) — maximum mean discrepancy
- Wang et al. (2022) — kernel projected Wasserstein
- Hu and Lin (2025) — max-sliced Wasserstein
- $\vdots$


Two-Sample Testing in $\mathbb{R}^d$

Works in $\mathbb{R}^d$:


- Gretton et al. (2012) — maximum mean discrepancy
- Wang et al. (2022) — kernel projected Wasserstein
- Hu and Lin (2025) — max-sliced Wasserstein
- $\vdots$

What about **beyond** $\mathbb{R}^d$? What if the observed samples are themselves **probability measures**?


Two-sample testing in $\mathcal{P}(\mathcal{P}(\mathbb{X}))$




Two-sample testing in $\mathcal{P}(\mathcal{P}(\mathbb{X}))$



Two-sample testing in $\mathcal{P}(\mathcal{P}(\mathbb{X}))$




Two-sample testing in $\mathcal{P}(\mathcal{P}(\mathbb{X}))$



Examples: income distributions, mortality proportions, temperature studies, ...

Two-sample testing in $\mathcal{P}(\mathcal{P}(\mathbb{X}))$




Examples: income distributions, mortality proportions, temperature studies, ...

→ Need to work at the level of **Random Probability Measures**

Hierarchical sample

(unknown) $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$

(latent) $\tilde{P}_1, \dots, \tilde{P}_n \stackrel{\text{i.i.d.}}{\sim} Q$




Hierarchical sample

(unknown) $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$

(latent) $\tilde{P}_1, \dots, \tilde{P}_n \stackrel{\text{i.i.d}}{\sim} Q$

Hierarchical Sample:

$$\mathbf{X}_{n,m} := \begin{pmatrix} X_{1,1} & X_{1,2} & \cdots & X_{1,m} & \tilde{P}_1 \stackrel{\text{i.i.d}}{\sim} \tilde{P}_1 \\ X_{2,1} & X_{2,2} & \cdots & X_{2,m} & \tilde{P}_2 \stackrel{\text{i.i.d}}{\sim} \tilde{P}_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ X_{n,1} & X_{n,2} & \cdots & X_{n,m} & \tilde{P}_n \stackrel{\text{i.i.d}}{\sim} \tilde{P}_n \end{pmatrix}$$



Hierarchical sample


(unknown) $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$

(latent) $\tilde{P}_1, \dots, \tilde{P}_n \stackrel{\text{i.i.d.}}{\sim} Q$

Hierarchical Sample:

$$\mathbf{X}_{n,m} := \begin{pmatrix} X_{1,1} & X_{1,2} & \cdots & X_{1,m} & \tilde{P}_1 \stackrel{\text{i.i.d.}}{\sim} \tilde{P}_1 \\ X_{2,1} & X_{2,2} & \cdots & X_{2,m} & \tilde{P}_2 \stackrel{\text{i.i.d.}}{\sim} \tilde{P}_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ X_{n,1} & X_{n,2} & \cdots & X_{n,m} & \tilde{P}_n \stackrel{\text{i.i.d.}}{\sim} \tilde{P}_n \end{pmatrix}$$

Notation: $\mathbf{X}_{n,m} \sim Q$



Hierarchical Samples in the Literature

- Baxter (2000) — meta-learning, distribution over tasks
- Szabó et al. (2016) — distribution regression
- Haviv et al. (2024) — generative modelling over distributions
- Warren et al. (2025) — principal curves in Wasserstein space

Limitations of the Hierarchical Sampling Framework

- Each row of $\mathbf{X}_{n,m}$ has the law

$$\mathbb{P}(X_{i,1} \in A_1, \dots, X_{i,m} \in A_m) = \int_{\mathcal{P}(\mathbb{X})} \prod_{j=1}^m P(A_j) Q(dP), \quad i = 1, \dots, n.$$

Limitations of the Hierarchical Sampling Framework

- Each row of $\mathbf{X}_{n,m}$ has the law

$$\mathbb{P}(X_{i,1} \in A_1, \dots, X_{i,m} \in A_m) = \int_{\mathcal{P}(\mathbb{X})} \prod_{j=1}^m P(A_j) Q(dP), \quad i = 1, \dots, n.$$

- $\implies (X_{i,1}, \dots, X_{i,m})$ is exchangeable.

Limitations of the Hierarchical Sampling Framework

- Each row of $\mathbf{X}_{n,m}$ has the law

$$\mathbb{P}(X_{i,1} \in A_1, \dots, X_{i,m} \in A_m) = \int_{\mathcal{P}(\mathbb{X})} \prod_{j=1}^m P(A_j) Q(dP), \quad i = 1, \dots, n.$$

- $\implies (X_{i,1}, \dots, X_{i,m})$ is exchangeable.
- **Limitation.** $Q^1 \neq Q^2$ does *not* imply $\mathbf{X}_{n,m} \stackrel{d}{\neq} \mathbf{Y}_{n,m}$.

Limitations of the Hierarchical Sampling Framework

- Each row of $\mathbf{X}_{n,m}$ has the law

$$\mathbb{P}(X_{i,1} \in A_1, \dots, X_{i,m} \in A_m) = \int_{\mathcal{P}(\mathbb{X})} \prod_{j=1}^m P(A_j) Q(dP), \quad i = 1, \dots, n.$$

- $\implies (X_{i,1}, \dots, X_{i,m})$ is exchangeable.
- Limitation.** $Q^1 \neq Q^2$ does *not* imply $\mathbf{X}_{n,m} \stackrel{d}{\neq} \mathbf{Y}_{n,m}$.
- But for m large:**

$$Q^1 = Q^2 \iff \mathbf{X}_{n,m} \stackrel{d}{=} \mathbf{Y}_{n,m}.$$

Distance-based Two Sample Test

- **Problem.** $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; test $H_0 : Q^1 = Q^2$ vs $H_1 : Q^1 \neq Q^2$.

Distance-based Two Sample Test

- **Problem.** $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; test $H_0 : Q^1 = Q^2$ vs $H_1 : Q^1 \neq Q^2$.
- **Strategy.** Given hierarchical samples $\mathbf{X}_{n,m} \sim Q^1$ and $\mathbf{Y}_{n,m} \sim Q^2$, reject H_0 if $d(\mathbf{X}_{n,m}, \mathbf{Y}_{n,m})$ is large.

Distance-based Two Sample Test

- **Problem.** $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; test $H_0 : Q^1 = Q^2$ vs $H_1 : Q^1 \neq Q^2$.
- **Strategy.** Given hierarchical samples $\mathbf{X}_{n,m} \sim Q^1$ and $\mathbf{Y}_{n,m} \sim Q^2$, reject H_0 if $d(\mathbf{X}_{n,m}, \mathbf{Y}_{n,m})$ is large.
- **Ingredients.** A distance d ; a threshold.

Hierarchical Empirical Estimator

- $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$.

Hierarchical Empirical Estimator

- $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$.

- Analogue of the empirical measure in the i.i.d. case: $\tilde{P}_1, \dots, \tilde{P}_n \stackrel{\text{i.i.d.}}{\sim} Q$

$$\frac{1}{n} \sum_{i=1}^n \delta_{\tilde{P}_i}.$$

- But $\tilde{P}_1, \dots, \tilde{P}_n$ are not observed.

Hierarchical Empirical Estimator

- $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$.
- Analogue of the empirical measure in the i.i.d. case: $\tilde{P}_1, \dots, \tilde{P}_n \stackrel{\text{i.i.d.}}{\sim} Q$

$$\frac{1}{n} \sum_{i=1}^n \delta_{\tilde{P}_i}.$$

- But $\tilde{P}_1, \dots, \tilde{P}_n$ are not observed.
- Estimate each $\tilde{P}_i$ from its samples: $\hat{P}_{i,m} = \frac{1}{m} \sum_{j=1}^m \delta_{X_{i,j}}$

Hierarchical Empirical Estimator

- $Q \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$.
- Analogue of the empirical measure in the i.i.d. case: $\tilde{P}_1, \dots, \tilde{P}_n \stackrel{\text{i.i.d.}}{\sim} Q$

$$\frac{1}{n} \sum_{i=1}^n \delta_{\tilde{P}_i}.$$

- But $\tilde{P}_1, \dots, \tilde{P}_n$ are not observed.
- Estimate each $\tilde{P}_i$ from its samples: $\hat{P}_{i,m} = \frac{1}{m} \sum_{j=1}^m \delta_{X_{i,j}}$
- **Hierarchical Empirical Estimator:**

$$\hat{Q}_{n,m} = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}}.$$

Distance based Two Sample Test

- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; test $H_0 : Q^1 = Q^2$ vs $H_1 : Q^1 \neq Q^2$.
- Observe hierarchical samples

$$\mathbf{X}_{n,m} \sim Q^1, \quad \mathbf{Y}_{n,m} \sim Q^2.$$

- Build hierarchical empirical estimators:

$$\hat{Q}_{n,m}^1 = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}^1}, \quad \hat{Q}_{n,m}^2 = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}^2}.$$

Distance based Two Sample Test

- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; test $H_0 : Q^1 = Q^2$ vs $H_1 : Q^1 \neq Q^2$.
- Observe hierarchical samples

$$\mathbf{X}_{n,m} \sim Q^1, \quad \mathbf{Y}_{n,m} \sim Q^2.$$

- Build hierarchical empirical estimators:

$$\hat{Q}_{n,m}^1 = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}^1}, \quad \hat{Q}_{n,m}^2 = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}^2}.$$

- **Reject H_0** if

$$d\left(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2\right) \text{ is large.}$$

Metrics on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$: Wasserstein over Wasserstein

- $(\mathbb{X}, d_{\mathbb{X}})$ is a Polish space with bounded diameter.
- **Step 1:** Using $d_{\mathbb{X}}$ on $\mathbb{X}$, define Wasserstein distance $\mathcal{W}$ on $\mathcal{P}(\mathbb{X})$: for $P^1, P^2 \in \mathcal{P}(\mathbb{X})$,

$$\mathcal{W}(P^1, P^2) = \inf_{(X^1, X^2) \in \Gamma(P^1, P^2)} \mathbb{E}[d_{\mathbb{X}}(X^1, X^2)] ,$$

where $\Gamma(P^1, P^2)$ are couplings of P^1 and P^2 on $\mathbb{X} \times \mathbb{X}$.

Metrics on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$: Wasserstein over Wasserstein

- $(\mathbb{X}, d_{\mathbb{X}})$ is a Polish space with bounded diameter.
- **Step 1:** Using $d_{\mathbb{X}}$ on $\mathbb{X}$, define Wasserstein distance $\mathcal{W}$ on $\mathcal{P}(\mathbb{X})$: for $P^1, P^2 \in \mathcal{P}(\mathbb{X})$,

$$\mathcal{W}(P^1, P^2) = \inf_{(X^1, X^2) \in \Gamma(P^1, P^2)} \mathbb{E}[d_{\mathbb{X}}(X^1, X^2)] ,$$

where $\Gamma(P^1, P^2)$ are couplings of P^1 and P^2 on $\mathbb{X} \times \mathbb{X}$.

- **Step 2:** Using $\mathcal{W}$ on $\mathcal{P}(\mathbb{X})$, for $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$,

$$\mathcal{W}_{\mathcal{W}}(Q^1, Q^2) = \inf_{(\tilde{P}^1, \tilde{P}^2) \sim \Gamma(Q^1, Q^2)} \mathbb{E}[\mathcal{W}(\tilde{P}^1, \tilde{P}^2)] .$$

$\mathcal{W}_{\mathcal{W}}$ is called **Wasserstein over Wasserstein (WoW)**.

Metrics on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$: Hierarchical IPM


- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; **idea**: compare their pushforward measures on $\mathbb{R}$.

Metrics on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$: Hierarchical IPM

- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; **idea**: compare their pushforward measures on $\mathbb{R}$.
- For $f : \mathbb{X} \rightarrow \mathbb{R}$, let $\hat{f} : \mathcal{P}(\mathbb{X}) \rightarrow \mathbb{R}$, $\hat{f}(P) = \int_{\mathbb{X}} f dP$.


Metrics on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$: Hierarchical IPM

- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; **idea**: compare their pushforward measures on $\mathbb{R}$.
- For $f : \mathbb{X} \rightarrow \mathbb{R}$, let $\hat{f} : \mathcal{P}(\mathbb{X}) \rightarrow \mathbb{R}$, $\hat{f}(P) = \int_{\mathbb{X}} f dP$.



Metrics on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$: Hierarchical IPM

- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; **idea**: compare their pushforward measures on $\mathbb{R}$.
- For $f : \mathbb{X} \rightarrow \mathbb{R}$, let $\hat{f} : \mathcal{P}(\mathbb{X}) \rightarrow \mathbb{R}$, $\hat{f}(P) = \int_{\mathbb{X}} f dP$.



- **Hierarchical Integral Probability Metric** (Catalano and Lavenant 2024):

$$d_{\mathcal{F}}(Q^1, Q^2) = \sup_{f \in \mathcal{F}} \mathcal{W}(Q^1 \circ \hat{f}^{-1}, Q^2 \circ \hat{f}^{-1}).$$

- $\mathcal{F} = \text{Lip}_1(\mathbb{X}, \mathbb{R})$: $d_{\mathcal{F}} = d_{\text{Lip}}$.

Properties of $\mathcal{W}_{\mathcal{W}}$ and d_{Lip}

- $\mathcal{W}_{\mathcal{W}}$ and d_{Lip} are **Integral Probability Metrics**: $\mathcal{I}_{\mathcal{G}}(\mu, \nu) = \sup_{g \in \mathcal{G}} |\mu(g) - \nu(g)|$.

Properties of $\mathcal{W}_{\mathcal{W}}$ and d_{Lip}

- $\mathcal{W}_{\mathcal{W}}$ and d_{Lip} are **Integral Probability Metrics**: $\mathcal{I}_{\mathcal{G}}(\mu, \nu) = \sup_{g \in \mathcal{G}} |\mu(g) - \nu(g)|$.

$$\mathcal{W}_{\mathcal{W}}(Q^1, Q^2) = \sup_{h \in \text{Lip}_1(\mathcal{P}(\mathbb{X}), \mathbb{R}, \mathcal{W})} \left| \int_{\mathcal{P}(\mathbb{X})} h dQ^1 - \int_{\mathcal{P}(\mathbb{X})} h dQ^2 \right|,$$

$$d_{\text{Lip}}(Q^1, Q^2) = \sup_{h \in \mathcal{D}} \left| \int_{\mathcal{P}(\mathbb{X})} h dQ^1 - \int_{\mathcal{P}(\mathbb{X})} h dQ^2 \right|,$$

where $\mathcal{D} = \{h : \mathcal{P}(\mathbb{X}) \rightarrow \mathbb{R} : h(P) = g(P(f)), g \in \text{Lip}_1(\mathbb{R}, \mathbb{R}), f \in \text{Lip}_1(\mathbb{X}, \mathbb{R})\}$.

Properties of $\mathcal{W}_{\mathcal{W}}$ and d_{Lip}

- $\mathcal{W}_{\mathcal{W}}$ and d_{Lip} are **Integral Probability Metrics**: $\mathcal{I}_{\mathcal{G}}(\mu, \nu) = \sup_{g \in \mathcal{G}} |\mu(g) - \nu(g)|$.

$$\mathcal{W}_{\mathcal{W}}(Q^1, Q^2) = \sup_{h \in \text{Lip}_1(\mathcal{P}(\mathbb{X}), \mathbb{R}, \mathcal{W})} \left| \int_{\mathcal{P}(\mathbb{X})} h dQ^1 - \int_{\mathcal{P}(\mathbb{X})} h dQ^2 \right|,$$

$$d_{\text{Lip}}(Q^1, Q^2) = \sup_{h \in \mathcal{D}} \left| \int_{\mathcal{P}(\mathbb{X})} h dQ^1 - \int_{\mathcal{P}(\mathbb{X})} h dQ^2 \right|,$$

where $\mathcal{D} = \{h : \mathcal{P}(\mathbb{X}) \rightarrow \mathbb{R} : h(P) = g(P(f)), g \in \text{Lip}_1(\mathbb{R}, \mathbb{R}), f \in \text{Lip}_1(\mathbb{X}, \mathbb{R})\}$.

- $\mathcal{W}_{\mathcal{W}}$ metrizes weak convergence on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$.
- d_{Lip} metrizes weak convergence on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$ when $\mathbb{X}$ is compact.

Properties of $\mathcal{W}_{\mathcal{W}}$ and d_{Lip}

- $\mathcal{W}_{\mathcal{W}}$ and d_{Lip} are **Integral Probability Metrics**: $\mathcal{I}_{\mathcal{G}}(\mu, \nu) = \sup_{g \in \mathcal{G}} |\mu(g) - \nu(g)|$.

$$\mathcal{W}_{\mathcal{W}}(Q^1, Q^2) = \sup_{h \in \text{Lip}_1(\mathcal{P}(\mathbb{X}), \mathbb{R}, \mathcal{W})} \left| \int_{\mathcal{P}(\mathbb{X})} h dQ^1 - \int_{\mathcal{P}(\mathbb{X})} h dQ^2 \right|,$$

$$d_{\text{Lip}}(Q^1, Q^2) = \sup_{h \in \mathcal{D}} \left| \int_{\mathcal{P}(\mathbb{X})} h dQ^1 - \int_{\mathcal{P}(\mathbb{X})} h dQ^2 \right|,$$

where $\mathcal{D} = \{h : \mathcal{P}(\mathbb{X}) \rightarrow \mathbb{R} : h(P) = g(P(f)), g \in \text{Lip}_1(\mathbb{R}, \mathbb{R}), f \in \text{Lip}_1(\mathbb{X}, \mathbb{R})\}$.

- $\mathcal{W}_{\mathcal{W}}$ metrizes weak convergence on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$.
- d_{Lip} metrizes weak convergence on $\mathcal{P}(\mathcal{P}(\mathbb{X}))$ when $\mathbb{X}$ is compact.
- $\mathbb{E}[\mathcal{W}_{\mathcal{W}}(\hat{Q}_{n,m}, Q)] \rightarrow 0$ **slower than** any polynomial rate.
- $\mathbb{E}[d_{\text{Lip}}(\hat{Q}_{n,m}, Q)] \rightarrow 0$ at **polynomial rate**.

(Catalano and Lavenant 2024)

Distance based Two Sample Test

- $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$; test $H_0 : Q^1 = Q^2$ vs $H_1 : Q^1 \neq Q^2$.
- Observe hierarchical samples

$$\mathbf{X}_{n,m} \sim Q^1, \quad \mathbf{Y}_{n,m} \sim Q^2.$$

- Build hierarchical empirical estimators:

$$\hat{Q}_{n,m}^1 = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}^1}, \quad \hat{Q}_{n,m}^2 = \frac{1}{n} \sum_{i=1}^n \delta_{\hat{P}_{i,m}^2}.$$

- **Reject H_0** if

$$d\left(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2\right) \text{ is large.}$$

Theoretical Threshold via Rademacher Complexity

- **Idea:** Threshold $r(n, m, \theta)$ from a concentration inequality.

Theoretical Threshold via Rademacher Complexity

- **Idea:** Threshold $r(n, m, \theta)$ from a concentration inequality.

Theorem

Let $\mathcal{G}$, $\mathcal{H}$ be uniformly bounded function classes generating IPMs on $\mathcal{P}(\mathbb{X})$ and $\mathcal{P}(\mathcal{P}(\mathbb{X}))$ respectively, with $\mathcal{H} \subseteq \text{Lip}_1(\mathcal{P}(\mathbb{X}), \mathbb{R}, \mathcal{I}_{\mathcal{G}})$. Then for all $\theta \in (0, 1)$, $n, m \geq 1$,

$$\mathbb{P}\left(\left|\mathcal{I}_{\mathcal{H}}(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2) - \mathcal{I}_{\mathcal{H}}(Q^1, Q^2)\right| \geq r(n, m, \theta)\right) \leq \theta,$$

where

$$r(n, m, \theta) = O\left(\mathcal{R}_n(\mathcal{H}) + \mathcal{R}_m(\mathcal{G}) + \frac{1}{\sqrt{n}}\right).$$

Theoretical Threshold via Rademacher Complexity

$\mathbb{X} \subset \mathbb{R}^d$ bounded.


WoW:

$$\begin{cases} \mathcal{O}\left(\frac{\log \log n}{\log n} + \frac{1}{\sqrt{n}} + \frac{1}{\sqrt{m}}\right), & d = 1 \\ \mathcal{O}\left(\frac{\log \log n}{\log n} + \frac{\log n}{\sqrt{n}} + \frac{\log m}{\sqrt{m}}\right), & d = 2 \\ \mathcal{O}\left(\frac{\log \log n}{\log n} + \frac{1}{n^{1/d}} + \frac{1}{m^{1/d}}\right), & d > 2. \end{cases}$$

HIPM:


$$\begin{cases} \mathcal{O}\left(\frac{1}{\sqrt{n}} + \frac{1}{\sqrt{m}}\right), & d = 1 \\ \mathcal{O}\left(\frac{\log n}{\sqrt{n}} + \frac{\log m}{\sqrt{m}}\right), & d = 2 \\ \mathcal{O}\left(\frac{1}{n^{1/d}} + \frac{1}{m^{1/d}}\right), & d > 2. \end{cases}$$

Practical Threshold via Permutation




Practical Threshold via Permutation

- Observe $d(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$. **Idea:** Estimate its $(1 - \theta)$ -quantile under H_0 .




Practical Threshold via Permutation

- Observe $d(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$. **Idea:** Estimate its $(1 - \theta)$ -quantile under H_0 .
- $\hat{P}_{1,m}^1, \dots, \hat{P}_{n,m}^1, \hat{P}_{1,m}^2, \dots, \hat{P}_{n,m}^2 \stackrel{\text{i.i.d}}{\sim}$ under H_0



Practical Threshold via Permutation


- Observe $d(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$. **Idea:** Estimate its $(1 - \theta)$ -quantile under H_0 .
- $\hat{P}_{1,m}^1, \dots, \hat{P}_{n,m}^1, \hat{P}_{1,m}^2, \dots, \hat{P}_{n,m}^2 \stackrel{\text{i.i.d.}}{\sim}$ under H_0
- For $s = 1, \dots, S$: compute $d(\hat{Q}_{n,m}^{1,(s)}, \hat{Q}_{n,m}^{2,(s)})$



Practical Threshold via Permutation

- Observe $d(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$. **Idea:** Estimate its $(1 - \theta)$ -quantile under H_0 .
- $\hat{P}_{1,m}^1, \dots, \hat{P}_{n,m}^1, \hat{P}_{1,m}^2, \dots, \hat{P}_{n,m}^2 \stackrel{\text{i.i.d.}}{\sim}$ under H_0
- For $s = 1, \dots, S$: compute $d(\hat{Q}_{n,m}^{1,(s)}, \hat{Q}_{n,m}^{2,(s)})$
- **Threshold:** $(1 - \theta)$ -quantile of

$$d(\hat{Q}_{n,m}^{1,(1)}, \hat{Q}_{n,m}^{2,(1)}), \dots, d(\hat{Q}_{n,m}^{1,(S)}, \hat{Q}_{n,m}^{2,(S)})$$



Practical Threshold via Permutation

- **Test statistic:** $T_{n,m} = \sqrt{\frac{n}{2}} d_{\text{Lip}}(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$
- **Permutation threshold:** for significance level $\theta \in (0, 1)$,

$$r(n, m, \theta) = \inf \left\{ t \in \mathbb{R} : \mathbb{P}(\tilde{T}_{n,m} > t \mid \mathbf{X}_{n,m}, \mathbf{Y}_{n,m}) \leq \theta \right\}.$$

Practical Threshold via Permutation

- **Test statistic:** $T_{n,m} = \sqrt{\frac{n}{2}} d_{\text{Lip}}(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$

- **Permutation threshold:** for significance level $\theta \in (0, 1)$,

$$r(n, m, \theta) = \inf \left\{ t \in \mathbb{R} : \mathbb{P}(\tilde{T}_{n,m} > t \mid \mathbf{X}_{n,m}, \mathbf{Y}_{n,m}) \leq \theta \right\}.$$

- **Testing scheme:** Reject H_0 if $T_{n,m} > r(n, m, \theta)$.

Practical Threshold via Permutation

- **Test statistic:** $T_{n,m} = \sqrt{\frac{n}{2}} d_{\text{Lip}}(\hat{Q}_{n,m}^1, \hat{Q}_{n,m}^2)$
- **Permutation threshold:** for significance level $\theta \in (0, 1)$,

$$r(n, m, \theta) = \inf \left\{ t \in \mathbb{R} : \mathbb{P}(\tilde{T}_{n,m} > t \mid \mathbf{X}_{n,m}, \mathbf{Y}_{n,m}) \leq \theta \right\}.$$

- **Testing scheme:** Reject H_0 if $T_{n,m} > r(n, m, \theta)$.

Theorem

Let $\mathbb{X} = [a, b]$ and $Q^1, Q^2 \in \mathcal{P}(\mathcal{P}(\mathbb{X}))$. Then for all $\theta \in (0, 1)$


- if $Q^1 = Q^2$: $\lim_{n,m \rightarrow \infty} \mathbb{P}(T_{n,m} > r(n, m, \theta)) \leq \theta$.
- if $Q^1 \neq Q^2$: $\lim_{n,m \rightarrow \infty} \mathbb{P}(T_{n,m} > r(n, m, \theta)) = 1$.

WoW vs. HIPM: Type I error and Power

$$n = m = n_{\text{perm}} = 100.$$

Type I error

$$\text{DP}(\alpha = 1, P_0 = \text{Uniform}(0, 1))$$




- Both methods **control Type I error**, including $\mathbb{X} = \mathbb{R}$ where theory does not guarantee it.

WoW vs. HIPM: Type I error and Power

$$n = m = n_{\text{perm}} = 100.$$

Type I error


$$\text{DP}(\alpha = 1, P_0 = \text{Uniform}(0, 1))$$



Power: A.

$$\text{DP}(\alpha = 1, P_0 = \text{Beta}(a, b)),$$

varying P_0




- Both methods **control Type I error**, including $\mathbb{X} = \mathbb{R}$ where theory does not guarantee it.
- **A:** HIPM slightly outperforms WoW.

WoW vs. HIPM: Type I error and Power

$$n = m = n_{\text{perm}} = 100.$$

Type I error


$$\text{DP}(\alpha = 1, P_0 = \text{Uniform}(0, 1))$$



Power: A.

$$\text{DP}(\alpha = 1, P_0 = \text{Beta}(a, b)),$$


varying P_0



Power: B.

$$\text{DP}(\alpha = 1, P_0 = \mathcal{N}(0, 1)),$$

varying α



- Both methods **control Type I error**, including $\mathbb{X} = \mathbb{R}$ where theory does not guarantee it.
- **A:** HIPM slightly outperforms WoW.
- **B:** WoW outperforms HIPM.

- **Fréchet test** (Dubey and Müller 2019)
- **Energy test** (Székely and Rizzo 2004)

Comparison with Existing Methods

Model: $\tilde{P}^i | \tilde{\mu} = \mathcal{N}(\tilde{\mu}, 1)$, $\tilde{\mu} \sim \mathcal{N}_T(\mu_{0i}, \sigma_{0i}^2, -10, 10)$.


$n = 100$, $m = 200$, $\theta = 0.05$.

Comparison with Existing Methods

Model: $\tilde{P}^i | \tilde{\mu} = \mathcal{N}(\tilde{\mu}, 1)$, $\tilde{\mu} \sim \mathcal{N}_T(\mu_{0i}, \sigma_{0i}^2, -10, 10)$.

$n = 100$, $m = 200$, $\theta = 0.05$.

Mean shift: $\sigma_{0i} = 0.5$, $\mu_{01} = 0$, $\mu_{02} = \delta$



- **Mean shift:** Energy, WoW, HIPM outperform Fréchet.



Comparison with Existing Methods

Model: $\tilde{P}^i | \tilde{\mu} = \mathcal{N}(\tilde{\mu}, 1)$, $\tilde{\mu} \sim \mathcal{N}_T(\mu_{0i}, \sigma_{0i}^2, -10, 10)$.

$n = 100$, $m = 200$, $\theta = 0.05$.

Mean shift: $\sigma_{0i} = 0.5$, $\mu_{01} = 0$, $\mu_{02} = \delta$

Variance shift: $\mu_{0i} = 0$, $\sigma_{01} = 0.2$, $\sigma_{02} = 0.2\tau$



- **Mean shift:** Energy, WoW, HIPM outperform Fréchet.
- **Variance shift:** Fréchet > WoW > HIPM > Energy.
- All methods control Type I error ($\delta = 0$, $\tau = 1$).

Comparison with Existing Methods

- $Q^1 \neq Q^2$ but **same** Fréchet mean and variance.

$\tilde{P}|\tilde{\mu} = \mathcal{N}(\tilde{\mu}, 1)$ in both cases, with

$$Q^1 : \tilde{\mu} \sim \frac{1}{2}\delta_{-1} + \frac{1}{2}\delta_1,$$

$$Q^2 : \tilde{\mu} \sim \frac{1}{8}\delta_{-2} + \frac{3}{4}\delta_0 + \frac{1}{8}\delta_2.$$

Test family: (Q^1, M_λ) , $M_\lambda = \lambda Q^1 + (1-\lambda)Q^2$,
 $\lambda \in [0, 1]$.

$\lambda = 1$: Type I error; $\lambda < 1$: power.

Comparison with Existing Methods

- $Q^1 \neq Q^2$ but **same** Fréchet mean and variance.

$\tilde{P}|\tilde{\mu} = \mathcal{N}(\tilde{\mu}, 1)$ in both cases, with


$$Q^1 : \tilde{\mu} \sim \frac{1}{2}\delta_{-1} + \frac{1}{2}\delta_1,$$

$$Q^2 : \tilde{\mu} \sim \frac{1}{8}\delta_{-2} + \frac{3}{4}\delta_0 + \frac{1}{8}\delta_2.$$

Test family: (Q^1, M_λ) , $M_\lambda = \lambda Q^1 + (1-\lambda)Q^2$,
 $\lambda \in [0, 1]$.

$\lambda = 1$: Type I error; $\lambda < 1$: power.

- Fréchet test cannot detect $Q^1 \neq M_\lambda$.
- HIPM, WoW, and Energy correctly detect the discrepancy.





Application: Human Mortality Data

Data: Human Mortality Database, 1960–2010.

age-at-death distribution over ages $\{0, 1, \dots, 110\}$ for each country and year.

- Group 1: Belarus, Bulgaria, Czechia, Estonia, Hungary, Latvia, Lithuania, Poland, Russia, Slovakia, Ukraine
- Group 2: Australia, Austria, Belgium, Canada, Denmark, Finland, France, Iceland, Ireland, Italy, Japan, Luxembourg, Netherlands, New Zealand, Norway, Spain, Sweden, Switzerland, UK, USA.

Application: Human Mortality Data



Quantile plots of age-at-death distributions for years 1960, 1992, 2008.

- The years capture the Cold War era, post-Soviet collapse, and a more recent period.
- Differences are visually apparent in later periods.

Application: Human Mortality Data


Methods:

- HIPM, WoW
- **Averaging:** $\bar{P}^k = \frac{1}{n_k} \sum_i \hat{P}_i^k$; $\mathcal{W}_1(\bar{P}^1, \bar{P}^2)$.
- **Pooling:** aggregate all deaths per group; permute at observation level.


$$\theta = 0.05, n_{\text{perm}} = 100$$

Application: Human Mortality Data



Males
P-values



Females
P-values





Application: Human Mortality Data



- HIPM, WoW, and Averaging show similar trends: significant differences after the 1970s.

Application: Human Mortality Data



- HIPM, WoW, and Averaging show similar trends: significant differences after the 1970s.
- **Pooling** fails to control Type I error in hierarchical sampling setting.

-  Baxter, Jonathan (2000). **“A Model of Inductive Bias Learning”**. In: *Journal of Artificial Intelligence Research* 12, pp. 149–198.
-  Catalano, Marta and Hugo Lavenant (2024). **“Hierarchical Integral Probability Metrics: A distance on random probability measures with low sample complexity”**. In: *Proceedings of the 41st International Conference on Machine Learning* 235, pp. 5841–5861.
-  Dubey, Paromita and Hans-Georg Müller (2019). **“Fréchet analysis of variance for random objects”**. In: *Biometrika* 106.4, pp. 803–821.
-  Gretton, Arthur et al. (2012). **“A Kernel Two-Sample Test”**. In: *Journal of Machine Learning Research* 13.25, pp. 723–773.
-  Haviv, Doron et al. (2024). **“Wasserstein flow matching: Generative modeling over families of distributions”**. In: *arXiv preprint arXiv:2411.00698*.

-  Hu, Xiaoyu and Zhenhua Lin (Apr. 2025). **“Two-sample distribution tests in high dimensions via max-sliced Wasserstein distance and bootstrapping”**. In: *Biometrika* 112.2, asaf001.
-  Szabó, Zoltán et al. (2016). **“Learning Theory for Distribution Regression”**. In: *Journal of Machine Learning Research* 17.152, pp. 1–40.
-  Székely, G. J. and M. L. Rizzo (2004). **“Testing for Equal Distributions in High Dimensions”**. In: *InterStat* 5, pp. 1249–1272.
-  Wang, Jie, Rui Gao, and Yao Xie (2022). **“Two-Sample Test with Kernel Projected Wasserstein Distance”**. In: *Proceedings of The 25th International Conference on Artificial Intelligence and Statistics* 151, pp. 8022–8055.
-  Warren, Andrew et al. (2025). **“Principal Curves in Metric Spaces and the Space of Probability Measures”**. In: *arXiv preprint arXiv:2505.04168*.

Thank you for your attention!